

Time Series Analysis

Md Mominul Islam

12/3/2020

1. (5 points) Plot the crime rate data vs the year.

Answer

```
## Loading required package: carData

##
## Attaching package: 'dplyr'

## The following object is masked from 'package:car':
##
##      recode

## The following objects are masked from 'package:stats':
##
##      filter, lag


## The following objects are masked from 'package:base':
##
##      intersect, setdiff, setequal, union

## Registered S3 method overwritten by 'quantmod':
##      method      from
##      as.zoo.data.frame zoo
```

```
Final_data <- as.data.frame(Final_data)

#Plotting Crime Rate Data vs The Year
plot(Final_data[,2],type="o",pch=19,cex=0.8,xlab='Year',
      xlim=c(1,23),xaxt= 'n', ylab='Crime Rate',
      main= "Crime Rate Data vs The Year")
axis(1, seq(1,22,1), Final_data[seq(1,22,1),1])
```

Crime Rate Data vs The Year




2. (10 points) Calculate and plot the sample autocorrelation function (ACF) and variogram. List the first 10 values of ACF and variogram respectively.

Answer


```
# Defining the variogram function
variogram_func <- function(x, lag) {
  x <- as.matrix(x) # Make sure the x is a vector. It represents the observations of y_t.
  Lag <- NULL
  var_k <- NULL
  vario <- NULL
  for (k in 1:lag) {
    Lag[k] <- k
    var_k[k] <- sd(diff(x, k))^2
    vario[k] <- var_k[k] / var_k[1]
  }
  return(as.data.frame(cbind(Lag, vario)))
}

#calculating ACF
ACF_Final_data <- acf(Final_data[,2], lag.max=25, type="correlation")
```



```
#calculating Variogram
x <- Final_data$Rate
lag_length <- length(x) /2
lag_cheese <- 1:lag_length
z <- variogram_func(x, lag_length)
variogram_Final_data <- z$vario
#Ploting Variogram
plot(variogram_Final_data, type="o",
      main="Variogram_original data", pch=19,cex=.8, col="black" )
```

Variogram_original data



```
#Table for ACF and Variogram
ACF <- ACF_Final_data$acf[1:10]
Variogram <- variogram_Final_data[1:10]
Variogram <- as.data.frame(Variogram)

Table <- data.frame (
  ACF,
  Variogram
)
kable(Table,
  caption = 'ACF and Variogram')
```

Table 1: ACF and Variogram

ACF	Variogram
1.0000000	1.000000
0.8984053	3.377662
0.7374591	6.703893
0.5532590	11.173629
0.3459439	16.429739
0.1236614	21.957853
-0.1018317	25.616645
-0.2854537	27.291365
-0.4133137	27.131428

ACF	Variogram
-0.4812716	25.404944

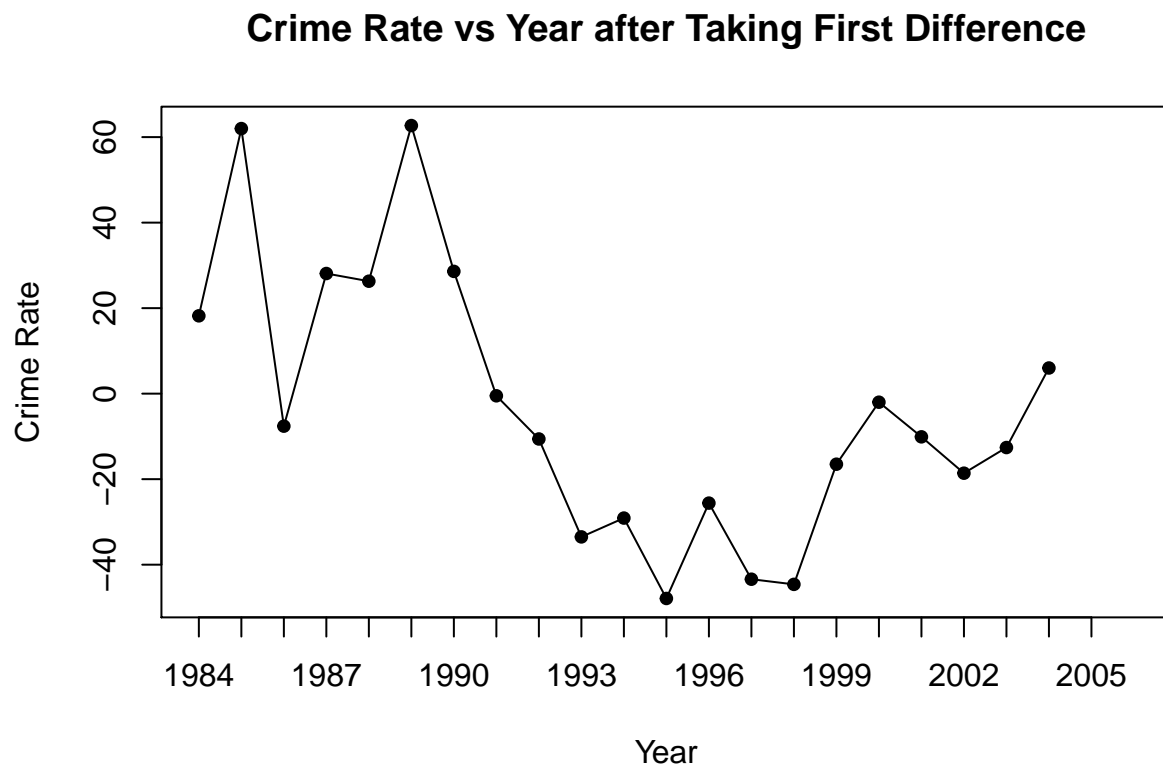
3. (5 points) Is there an indication of nonstationary behavior in the time series? Why or why not?

Answer

4. (10 points) Calculate and plot the first difference of the timeseries. Show the first 10 differences.

Answer

```
#Calculating First Difference Data
Final_data.df1 <- diff(Final_data[,2])
#Plotting First Difference Data
plot(Final_data.df1, type="o", pch=19, cex=0.8, xlab='Year',
      xlim=c(1,23), xaxt= 'n', ylab='Crime Rate',
      main= "Crime Rate vs Year after Taking First Difference")
axis(1, seq(1,22,1), Final_data[seq(1,22,1),1])
```



```
#Showing First 10 Differences
Final_data_diff <- cbind(Final_data[1:10,1], Final_data[1:10,2], Final_data.df1[1:10])
Final_data_diff <- as.data.frame(Final_data_diff)
#Replacing with suitable one
```

```
colnames(Final_data_diff) = c("Year", "Crime Rate", "Difference Data")
kable(Final_data_diff,
      caption = 'First Differences of the Time Series')
```


Table 2: First Differences of the Time Series

Year	Crime Rate	Difference Data
1984	539.9	18.2
1985	558.1	62.0
1986	620.1	-7.6
1987	612.5	28.1
1988	640.6	26.3
1989	666.9	62.7
1990	729.6	28.6
1991	758.2	-0.5
1992	757.7	-10.6
1993	747.1	-33.5

5. (10 points) Compute the sample autocorrelation function (ACF) and variogram of the first differences.

Answer

```
#calculating ACF
ACF_Final_data_dif1 <-acf(Final_data.df1, lag.max=25,type="correlation")
```

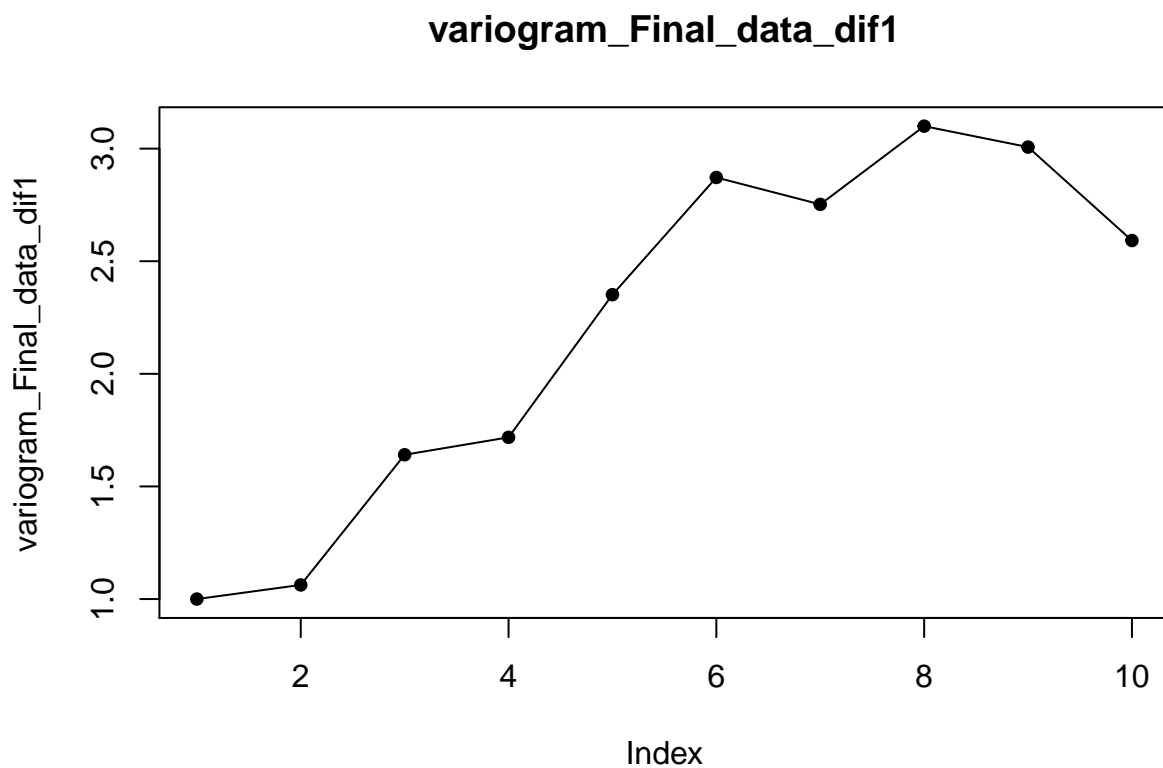


```

#calculating Variogram
x <- Final_data.df1
lag_length <- length(x) /2
lag_cheese <- 1:lag_length
z <- variogram_func(x, lag_length)
variogram_Final_data_dif1 <- z$vario

#Table for ACF and Variogram
ACF_dif1 <- ACF_Final_data_dif1$acf[1:10]
Variogram_dif1 <- variogram_Final_data_dif1[1:10]
Variogram_dif1 <- as.data.frame(Variogram_dif1)
plot(variogram_Final_data_dif1, type="o",
      main="variogram_Final_data_dif1", pch=19,cex=.8, col="black" )

```



```

Table_dif1 <- data.frame (
  ACF_dif1,
  Variogram_dif1
)
kable(Table_dif1,
      caption = 'ACF with First Differences',
      col.names = c('ACF D1', 'Variogram D1'))

```

Table 3: ACF with First Differences

ACF D1	Variogram D1
1.0000000	1.000000
0.6193703	1.062543
0.4985524	1.640599
0.3198224	1.717923
0.2879014	2.351421
0.1007391	2.871869
-0.1641531	2.752315
-0.2267162	3.099729
-0.3636319	3.006877
-0.3412978	2.591916

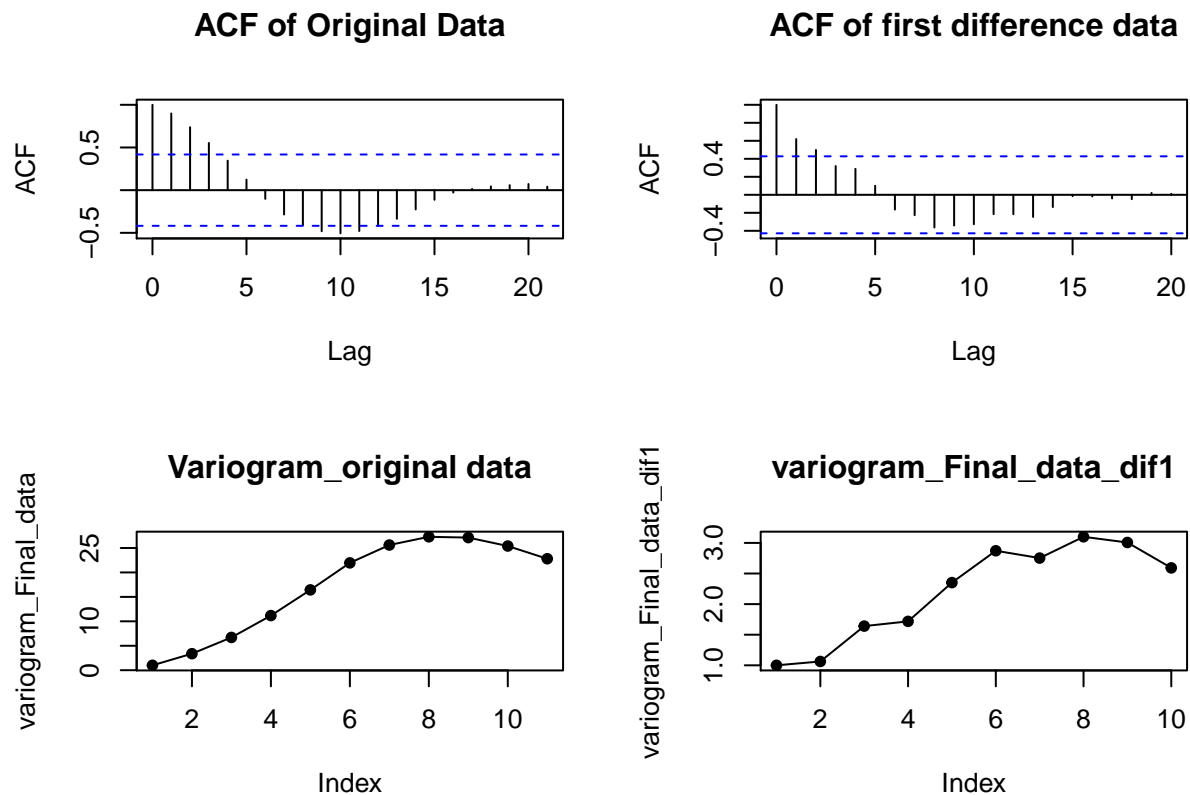
6. (5 points) What impact has differencing had on the time series?

Answer

```
#Sample autocorrelation function of original data
par(mfrow=c(2, 2))
ACF_Final_data <-acf(Final_data[,2],
                     lag.max=25,type="correlation",
                     main="ACF of Original Data")
#Sample autocorrelation function of first difference data
ACF_Final_data_dif1 <-acf(Final_data.df1, lag.max=25,type="correlation",
main="ACF of first difference data")

#Variogram of Original data
plot(variogram_Final_data, type="o",
     main="Variogram_original data",
     pch=19,cex=.8, col="black" )

#Variogram of First Difference Data
plot(variogram_Final_data_dif1, type="o", main="variogram_Final_data_dif1", pch=19,cex=.8, col="black" )
```

7. Develop an appropriate exponential smoothing forecasting procedure for the firstdifferencing data by answer the questions below.

```
#Defining Smoothing Function
firstsmooth <- function(y, lambda, start=y[1]) {
  ## here the initial value is set to the first y value
  ytilde <- y
  ytilde[1] <- lambda*y[1]+(1-lambda)*start
  for (i in 2:length(y)) {
    ytilde[i] <- lambda*y[i] + (1-lambda)*ytilde[i-1]
  }
  ytilde
}

## Defining Trigg leach smooth Function
tlsmooth<-function(y,gamma,y.tilde.start=y[1],lambda.start=1){
  T<-length(y)
  #Initialize the vectors
  Qt<-vector()
  Dt<-vector()
  y.tilde<-vector()
  lambda<-vector()
  err<-vector()
  #Set the starting values for the vectors
  lambda[1]=lambda.start
  y.tilde[1]=y.tilde.start
  Qt[1]<-0
```

```

Dt[1]<-0
err[1]<-0
for (i in 2:T){
  err[i]<-y[i]-y.tilde[i-1]
  Qt[i]<-gamma*err[i]+(1-gamma)*Qt[i-1]
  Dt[i]<-gamma*abs(err[i])+(1-gamma)*Dt[i-1]
  lambda[i]<-abs(Qt[i]/Dt[i])
  y.tilde[i]=lambda[i]*y[i] + (1-lambda[i])*y.tilde[i-1]
}
return(cbind(y.tilde,lambda,err,Qt,Dt))
}

##Function for measures of accuracy
measacc.fs<- function(y,lambda){
  out<- firstsmooth(y,lambda)
  T<-length(y)
  #Smoothed version of the original is the one step ahead prediction
  #Hence the predictions (forecasts) are given as
  pred<-c(y[1],out[1:(T-1)])
  prederr<- (y - pred)
  SSE<-sum(prederr*prederr)
  MAPE<-100*sum(abs(prederr/y))/T
  MAD<-sum(abs(prederr))/T
  MSD<-sum(prederr*prederr)/T
  ret1<-c(SSE,MAPE,MAD,MSD)
  names(ret1)<-c("SSE", "MAPE", "MAD", "MSD")
  return(ret1)
}

```

a. (10 points) Assume the first-difference data is a constant process. For R user, use the HoltWinters() function to find the optimum value of Lamda to smooth the data. For JMP user, specify the Lamda given by the software.

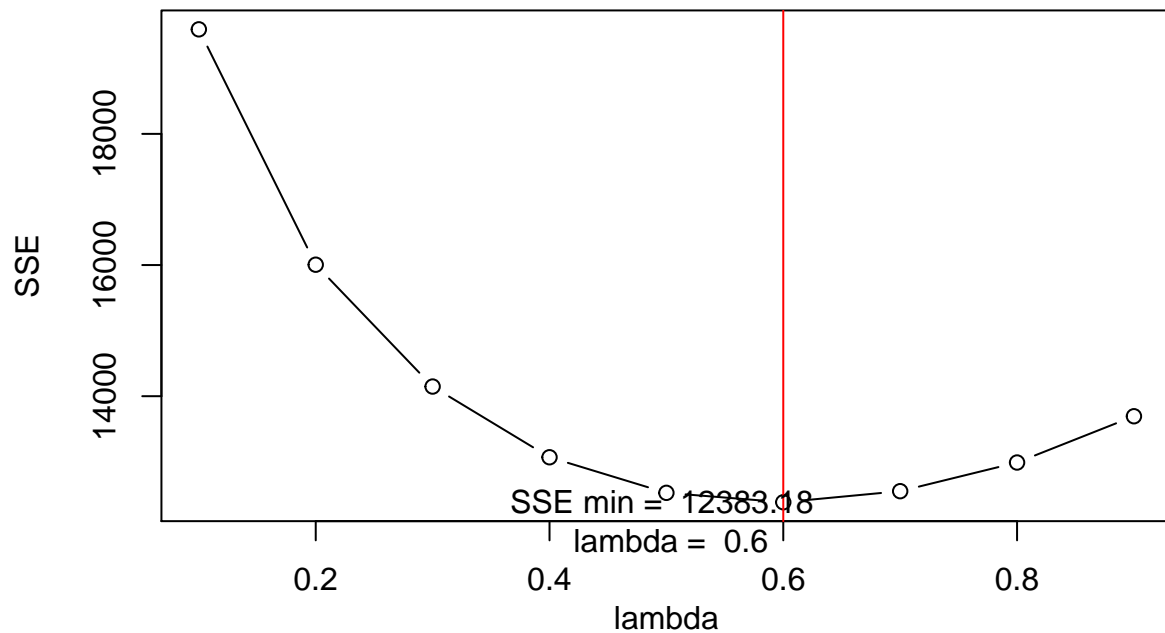
Answer

```

##Finding Out suitable Lamda Value for Low SSE
lambda.vec<-seq(0.1, 0.9, 0.1)
sse.mydata<-function(sc){measacc.fs(Final_data.df1,sc)[1]}
sse.vec<-sapply(lambda.vec, sse.mydata)
opt.lambda<-lambda.vec[sse.vec == min(sse.vec)]
plot(lambda.vec, sse.vec, type="b", main = "SSE vs. lambda\n",
xlab='lambda\n',ylab='SSE')
abline(v=opt.lambda, col = 'red')
mtext(text = paste("SSE min = ",
                    round(min(sse.vec),2), "\n lambda = ",
                    opt.lambda), side =1)

```

SSE vs. lambda



#Fitted Data

```
fit1 <- HoltWinters(Final_data.df1,alpha=0.3,beta=FALSE, gamma=FALSE)
fit1
```

```
## Holt-Winters exponential smoothing without trend and without seasonal component.
##
## Call:
## HoltWinters(x = Final_data.df1, alpha = 0.3, beta = FALSE, gamma = FALSE)
##
## Smoothing parameters:
##  alpha: 0.3
##  beta : FALSE
##  gamma: FALSE
##
## Coefficients:
##      [,1]
## a -9.529966
```

b. (10 points) Show the fitted values and corresponding SSE by using the lamda obtained in part a.

```
Final_Data_fitted<-firstsmooth(y=Final_data.df1,lambda=0.6)
data.frame(Final_Data_fitted)
```

```
##      Final_Data_fitted
```

```
## 1      18.2000000
## 2      44.4800000
## 3      13.2320000
## 4      22.1528000
## 5      24.6411200
## 6      47.4764480
## 7      36.1505792
## 8      14.1602317
## 9      -0.6959073
## 10     -20.3783629
## 11     -25.6113452
## 12     -38.9845381
## 13     -30.9538152
## 14     -38.4215261
## 15     -42.1286104
## 16     -26.7514442
## 17     -11.9005777
## 18     -10.8202311
## 19     -15.4880924
## 20     -13.7552370
## 21     -1.9020948
```

```
Final_Data_fitted.sse<- measacc.fs(Final_Data_fitted,0.6)
kable(Final_Data_fitted.sse)
```

	x
SSE	5237.82056
MAPE	245.08845
MAD	12.36447
MSD	249.42003

c. (5 points) Plot the fitted values and original values in a same plot.

```
#First Difference Data
```


```
Final_data_diff_all <- cbind(Final_data[,1],Final_data[,2], Final_data.df1)
```

```
## Warning in cbind(Final_data[, 1], Final_data[, 2], Final_data.df1): number of
## rows of result is not a multiple of vector length (arg 3)
```

```
Final_data_diff_all <-as.data.frame(Final_data_diff_all)
Final_data_diff_all$V2 <- NULL
```

```
plot(Final_data_diff_all[1:T,2], type='p', pch=16, cex=1,xlim=c(1,23),xaxt='n',ylim=c(-100,200),
     xlab='Year',xaxt= 'n', ylab='Crime Rate',
     main= "Forecasted Crime Rate Data vs The Year")
axis(1, seq(1,23,1), Final_data_diff_all[seq(1,23,1),1])
points(1:21,Final_Data_fitted, lty =2, col= "blue")
```

Forcasted Crime Rate Data vs The Year



d. (5 points) Assume the first-difference data shows a trend. Calculate the SSE. You can get it from the `HoltWinters()` function. Then compare the SSE with that of obtained in part b. What can you tell from the comparison?

```
fit1_trending <- HoltWinters(Final_data.df1,alpha=0.3,beta=0.3, gamma=FALSE)
fit1_trending$fitted
```

```
## Time Series:
## Start = 3
## End = 21
## Frequency = 1
##      xhat      level      trend
## 3 105.800000 62.000000 43.800000
## 4 105.374000 71.780000 33.594000
## 5 108.831140 82.191800 26.639340
## 6 103.283335 84.071798 19.2115374
## 7 106.667372 91.108335 15.5590372
## 8 91.780134 83.247160 8.5329737
## 9 64.323856 64.096094 0.2277617
## 10 35.331314 41.846699 -6.5153853
## 11 1.971716 14.681919 -12.7102036
## 12 -22.856457 -7.349799 -15.5066580
## 13 -48.130097 -30.369520 -17.7605769
## 14 -57.103936 -41.371068 -15.7328682
## 15 -67.492269 -52.992755 -14.4995139
## 16 -73.063798 -60.624588 -12.4392097
```

```
## 17 -63.443127 -56.094659 -7.3484679
## 18 -46.828775 -45.010189 -1.8185865
## 19 -34.323139 -35.810143 1.4870032
## 20 -26.704112 -29.606198 2.9020858
## 21 -18.301422 -22.472878 4.1714558
```

```
fit1_trending$SSE
```

```
## [1] 64423.21
```

e. (5 points) Suppose the first-difference is a constant process. Give the forecasts of the crime rate for years from 2006 to 2010.


Answer

```
# Forecast Values
lcpi<-0.6
cpi.smooth1<-firstsmooth(y=Final_data_diff_all[1:22,2],lambda=lcpi)
cpi.smooth2<-firstsmooth(y=cpi.smooth1,lambda=lcpi)
cpi.hat<-2*cpi.smooth1-cpi.smooth2
tau<-1:6
T<-length(cpi.smooth1)
cpi.forecast<-(2+tau*(lcpi/(1-lcpi)))*cpi.smooth1[T]-(1+tau*(lcpi/(1-lcpi)))*cpi.smooth2[T]
ctau<-sqrt(1+(lcpi/((2-lcpi)^3))*(10-14*lcpi+5*(lcpi^2)+2*
                                                    tau*lcpi*(4-3*lcpi)+2*(tau^2)*(lcpi^2)))

alpha.lev<-0.05
sig.est<- sqrt(var(Final_data_diff_all[2:22,2]- cpi.hat[1:21]))
cl<-qnorm(1-alpha.lev/2)*(ctau/ctau[1])*sig.est

#plot forecast model with prediction intervals
plot(Final_data_diff_all[1:T,2], type='o', pch=16, cex=1,xlim=c(1,30),xaxt='n',ylim=c(-100,200),
     xlab='Year',xaxt='n', ylab='Crime Rate',
     main= "Forecasted Crime Rate Data vs The Year")
axis(1, seq(1,29,1), Final_data_diff_all[seq(1,29,1),1])
lines(23:28,cpi.forecast, lty =2, col= "blue")
lines(23:28,cpi.forecast+cl, lty =4, col = "red")
lines(23:28,cpi.forecast-cl, lty =6, col = "green")
legend( x="topleft",
       legend=c("Forecast", "95% UPL", "95% LPL"),
       col=c("blue", "red", "green"), lwd=.1, lty=c(2,4,6,NA),
       pch=c(NA,NA,NA,15), merge=FALSE, cex= 1)
```

Forcasted Crime Rate Data vs The Year



8. a. (10 points) Develop an appropriate ARIMA model and a procedure for forecasting for the crime rate data. Specify the model and estimated parameters in the model. Hint: You can use the `auto.arima()` and `forecast()` functions to answer this question.

Answer

```
#Using auto.arima function
```

```
Final_Dataset_AutoArimaModel<-auto.arima(Final_data[,2])
```

```
Final_Dataset_AutoArimaModel
```

```
## Series: Final_data[, 2]
```

```
## ARIMA(1,2,0)
```

```
##
```

```
## Coefficients:
```

```
##      ar1
```

```
##      -0.4533
```

```
## s.e.   0.2091
```

```
##
```

```
## sigma^2 estimated as 623.5: log likelihood=-92.33
```

```
## AIC=188.67  AICc=189.37  BIC=190.66
```

```
Final_Dataset_AutoArimaModel_Forecast<-as.array(forecast(Final_Dataset_AutoArimaModel))
```

```
kable(Final_Dataset_AutoArimaModel_Forecast)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
23	466.7685	434.76854	498.7685	417.82878	515.7082
24	468.1591	409.22121	527.0969	378.02140	558.2967
25	467.8171	373.55509	562.0790	323.65581	611.9783
26	468.2604	334.67715	601.8437	263.96243	672.5584
27	468.3478	290.70970	645.9859	196.67377	740.0218
28	468.5965	243.08759	694.1055	123.71037	813.4827
29	468.7721	191.67339	745.8709	44.98618	892.5581
30	468.9809	136.90442	801.0573	-38.88624	976.8480
31	469.1746	78.89640	859.4528	-127.70440	1066.0536
32	469.3751	17.85539	920.8949	-221.16472	1159.9150

b. (5 points) Compare the AIC obtained from part a with that of obtained from ARIMA(0,1,0) model. Which model has a smaller AIC? What can you tell by this comparison?

Answer

```
# AR(1) model create with code ARIMA (0,1,0)
Final_data.arima<-arima(Final_data[,2],order=c(0, 1, 0))
Final_data.arima

##
## Call:
## arima(x = Final_data[, 2], order = c(0, 1, 0))
##
##
## sigma^2 estimated as 966:  log likelihood = -101.97,  aic = 205.93
```

c. (5 points) Show the 1- to 5- step ahead forecasts and corresponding 95% prediction intervals for the crime rate. Show only the results/outputs. Calculation process or formula are not required.

Answer

```
plot(Final_data[,2],type="o",pch=19,cex=0.8,
     xlab='Year',xaxt='n', ylab='Crime Rate',
     xlim=c(1,28),xaxt='n', ylim=c(0,1000),
     main= "Forecasted Crime Rate Data vs The Year")
axis(1, seq(1,27,1), Final_data[seq(1,27,1),1])
lines(Final_Dataset_AutoArimaModel_Forecast$mean,col="blue", lty= 2)
lines(Final_Dataset_AutoArimaModel_Forecast$upper[,2], col="red", lty= 4)
lines(Final_Dataset_AutoArimaModel_Forecast$lower[,2], col="green",lty= 6)
legend( x="bottomleft",
       legend=c("Forecast", "95% UPL", "95% LPL"),
       col=c("blue", "red", "green"), lwd=.1, lty=c(2,4,6,NA),
       pch=c(NA,NA,NA,15), merge=FALSE, cex= .7 )
```


Forecasted Crime Rate Data vs The Year

